

Copula-Based Anomaly Detection in Multivariate Time Series

Anomaly detection seeks unusual or unexpected events in data. In **multivariate time series**, this task is especially challenging because many variables interact. Joint anomalies may only appear when variables co-vary unusually, so methods must account for their dependencies ¹. For example, extreme climate events or system faults often involve simultaneous deviations across multiple sensors. A suitable detector thus models the joint distribution of the series and flags points in the joint tails. Unlike univariate outliers (single-sensor spikes), multivariate anomalies can manifest as subtle deviations from normal **patterns** of co-movement. Pattern detection (e.g. learning repeated motifs or sub-sequences) provides a baseline of “normal” behavior, and any subsequence that significantly **deviates** from these learned patterns is anomalous ². In practice, one often models the typical data patterns (via clustering, Gaussian models, etc.) and then identifies points whose joint behavior falls outside this model. For instance, Verma et al. define a multivariate sequence as anomalous whenever any subsequence departs significantly from the normal pattern of a Gaussian mixture model ². Thus, anomaly detection and pattern mining are two sides of the same coin: anomalies are those data patterns that are incompatible with the normal patterns detected in the data.

Copula Theory for Modeling Dependence

A **copula** is a function that couples multivariate joint distributions to their one-dimensional marginals. By Sklar’s theorem ³, any multivariate distribution $F(x_1, \dots, x_d)$ with continuous marginals $F_i(x_i)$ can be written as

$$F(x_1, \dots, x_d) = C(F_1(x_1), \dots, F_d(x_d)),$$

where $C : [0, 1]^d \rightarrow [0, 1]$ is a copula function. Intuitively, a copula describes the *dependence* among variables after transforming each marginal to a uniform scale. This separates marginal modeling (e.g. fitting the univariate distributions of each series) from dependence modeling. In practice one can freely choose a copula family for the joint and independent marginal distributions; this modular approach is a major advantage ³ ⁴.

Different copula families capture different dependence features. **Elliptical copulas** (e.g. Gaussian, Student- t) are symmetric: a Gaussian copula encodes linear correlation but has no tail dependence, while a Student- t copula adds symmetric heavy tails. In contrast, **Archimedean copulas** (one-parameter families like Clayton, Gumbel, Frank, etc.) can model asymmetric dependence. For example, the Clayton copula emphasizes **lower-tail** co-movement (joint extremes at the low end), whereas the Gumbel copula emphasizes **upper-tail** co-movement ⁵. The Frank copula has no tail dependence. A useful summary from Nagler & Czado (2021) is: “Gaussian and Frank copulas do not exhibit tail dependence ($\lambda=0$), the Student- t has symmetric tail dependence, while Clayton and Gumbel have only lower or upper tail dependence, respectively.” ⁵. Figure 5 (below) illustrates these behaviors for a toy bivariate example. The Gaussian and Frank copulas yield light, symmetric contours, whereas the t-copula shows much higher joint probability in

the corner (symmetric heavy tails). The Gumbel case shows high probability in the top-right corner (upper-tail), and Clayton in the bottom-left (lower-tail).




Figure: Joint PDF, CDF and contour plots for five common bivariate copulas (Gaussian, t , Gumbel, Clayton, Frank). The tail-dependence patterns differ: e.g. (a) Gaussian has no joint tail mass, (b) t has heavy symmetric tails, (c) Gumbel emphasizes the upper-right tail, (d) Clayton the lower-left. This flexibility lets copulas model extreme co-movement beyond linear correlation ⁵.

Mathematically, an d -dimensional copula $C(u_1, \dots, u_d)$ is a joint CDF with uniform marginals. In practice, choosing an appropriate copula family and estimating its parameters (via maximum likelihood or inversion of rank correlations) allows one to capture complex dependencies in the data. **Vine copulas** (pair-copula constructions) generalize this further: they decompose a high-dimensional copula into a cascade of bivariate copulas arranged in a tree structure ⁶. This overcomes the limitations of single-form copulas in high dimensions. In summary, copulas provide a rich toolkit for modeling multivariate dependencies and tail behavior, which is crucial for anomaly detection in multi-sensor or high-dimensional time series.

Joint Anomaly Detection via Pooled p -Values and Copulas

To detect anomalies jointly across many series, one strategy is to first obtain an anomaly *score* or *p-value* from each individual series at each time, and then **combine** these signals. Standard methods like Fisher's or Stouffer's tests combine independent p -values, but these break down when the series are correlated. Copulas offer a principled way to aggregate dependent p -values. One can treat the marginal anomaly p -values u_i from each series as approximately $\text{Uniform}(0,1)$, then fit a multivariate copula to the vector $(u_1, \dots, u_d)$. The copula captures the dependence of significance across series. We can then compute a global p -value for the event that all (or some) series are simultaneously extreme.

A notable implementation of this idea is **COPOD** (Copula-based Outlier Detection) by Zhao et al. (2020) ⁷. COPOD first builds the empirical copula of the data and then computes, for each observation, the joint tail probability under this copula. Intuitively, this gives an “anomalous p -value” for each point ⁷. A small tail

probability means the point is extreme in the joint distribution. Crucially, COPOD makes no parametric assumptions and is parameter-free and interpretable ⁷ . More generally, there is a rich statistical literature on using copulas in multiple hypothesis testing to model dependence among p-values. For example, Bayesian approaches have used multivariate copulas to combine dependent test statistics in genomics ⁸ . In the anomaly context, this translates to: *fit a copula to the vector of per-series p-values; then the copula's joint CDF gives the probability that the point is as or more extreme than observed*. This yields a pooled p-value that correctly accounts for inter-series correlation.

Historical Background and Research Developments

Copula theory originated in probability in 1959 (Sklar’s theorem) ³ and gained prominence in finance (e.g. the Gaussian copula for CDOs). The early 2000s saw rapid development of copula models (Archimedean families, pair-copula vines, etc.) for multivariate modeling ⁶ . In the context of time series, Aas *et al.* (2009) and Brechmann & Czado (2013) pioneered using copulas (especially vines) to model financial time series and capture risk. Czado & Nagler (2014) proposed copula autoregression (COPAR) models to handle both serial and cross-sectional dependence. These efforts laid the groundwork for copulas in dynamic data.

Copula methods were adopted for anomaly and pattern detection only more recently. A key development was **Zhao *et al.* (2020)** introducing COPOD ⁷ , a fast, non-parametric algorithm for multivariate outlier detection using empirical copulas. This sparked attention in the anomaly detection community (it is included in Python’s PyOD toolkit, for example). In parallel, Horváth *et al.* (2019) proposed a **copula-based anomaly scoring** method for large-scale systems ⁹ . They modeled the joint distribution of high-dimensional network performance counters with a copula, explicitly targeting fat-tailed behavior: “Since rare events are present in the tails of the distribution, we use copula functions that are able to model the fat-tailed distributions well” ⁹ . Their method also localized which features drove the anomaly, supporting root-cause analysis.

Another line of work is **subspace scanning via copulas**. Dorigo & Ferroni (2021) introduced *RanBox*, which systematically scans copula-transformed data subspaces for anomalous “clusters” by minimizing a uniformity p-value ¹⁰ . In effect, RanBox uses copulas to define a standardized space and searches for local overdensities that indicate anomalies. This has been applied in high-energy physics (e.g., finding rare events in particle data).

More recently, copulas have been applied in domain-specific anomaly problems. For instance, Tenconi *et al.* (2024) used copula models for fraud detection in finance ¹¹ ¹² . They demonstrated that a copula-Gaussian GLM captured complex dependencies among transaction features, yielding better fraud probabilities than traditional methods. Other examples include copula-based graphical models for network intrusion classification, and statistical control charts using copula to monitor industrial processes. Table 1 (below) summarizes representative milestones.

Year	Contribution	Domain/Topic
1959	Sklar’s Theorem ³	Theory of copulas
2009	Pair-copula (vine) models (Aas <i>et al.</i>)	Financial time series modeling
2013	Vine copulas for risk (Brechmann & Czado)	Portfolio risk modeling

Year	Contribution	Domain/Topic
2014	Copula autoregression (Czado & Nagler)	Multivariate time series models
2019	Copula anomaly scoring (Horváth <i>et al.</i>) ⁹	Network monitoring
2020	COPOD outlier detection (Zhao <i>et al.</i>) ⁷	Unsupervised outlier detection
2021	RanBox copula scan (Dorigo & Ferroni) ¹⁰	High-energy physics (local outliers)
2024	Copula-GLM fraud detection (Tenconi <i>et al.</i>) ¹¹	Financial forensics

Comparison with Classical, ML and Deep Methods

Common multivariate anomaly detection techniques include classical statistics, machine learning, and deep learning. Classical methods (Hotelling's T^2 , multivariate CUSUM) assume Gaussian models and monitor covariance deviations ¹³. For example, Hotelling's T^2 essentially measures the Mahalanobis distance from a mean; anomalies are flagged when T^2 exceeds a threshold ¹³. These work under normality assumptions but **miss anomalies in the tails or nonlinear interactions**. PCA-based methods reduce dimensionality and detect outliers by reconstruction error, but they capture only linear subspaces. In one case study, a linear PCA-canonical-correlation approach **overestimated anomalies**, whereas a copula-based method remained robust when anomalies occurred in individual variables ¹⁴.

Unsupervised machine-learning methods such as k-NN / LOF, Isolation Forest, or one-class SVM make fewer distributional assumptions. They can handle arbitrary data but typically **ignore the specific dependence structure**: for instance, Isolation Forest builds random trees that treat each dimension independently. These methods often require careful tuning and may produce false alarms if features correlate unexpectedly. Deep learning approaches (autoencoders, LSTM/GRU networks, Transformers, graph neural nets) have become popular for multivariate time series ¹⁵. They can model complex temporal and cross-variable patterns, and have advanced anomaly detection in domains like IoT and finance ¹⁵. However, deep models are data-hungry, sensitive to hyperparameters, and usually opaque (hard to interpret or explain anomalies).

Copula-based methods **complement** these approaches by explicitly modeling the joint distribution. They make minimal parametric assumptions (empirical copulas are nonparametric) yet capture tail behavior that linear models miss ⁵. A copula detector can provide an interpretable "joint p-value" score for each observation, directly flagging extreme joint events. Unlike PCA, copulas do not suffer from linear model bias in the presence of anomalies ¹⁴. Compared to generic ML methods, copulas incorporate known statistical principles and often require fewer tuning parameters. In practice, one may even combine techniques: for instance, use PCA to reduce dimensionality and then apply a copula to model the remaining joint distribution. Table 2 below summarizes key differences.

- **Classical (e.g. Hotelling's T^2):** Assumes multivariate normality; uses covariance to detect outliers ¹³. Fails under skewed or heavy-tailed data; does not target tail anomalies.
- **PCA/Subspace methods:** Capture linear correlations and identify points outside the principal components. Easy but misses nonlinear dependencies; can be overwhelmed by a single anomalous feature (linear methods overestimate anomalies) ¹⁴.

- **Machine Learning (LOF, IF, SVM):** Nonparametric detectors that flag low-density points. Flexible but often treat features independently or rely on distance metrics; can struggle when anomalies affect correlation structure.
- **Deep Learning (LSTM, Autoencoder, VAE, GAN, Transformer)** ¹⁵ : Learn complex spatio-temporal patterns; excel with large labeled/unlabeled data. Powerful for sequential anomalies but require extensive training and lack transparency.
- **Copula-based methods:** Explicitly model joint CDF; detect anomalies as joint tail events. Capture non-linear and tail dependencies ⁵ , yielding interpretable probability scores. Complement other methods by accounting for dependence. Copula methods may be computationally heavier in high dimensions (though vines help ⁶) and require enough data to estimate the copula accurately.

Case Studies and Applications

Copula-based anomaly detection has been successfully applied in diverse real-world domains:

- **Finance and Fraud Detection:** In financial data (stock prices, transaction records), anomalies often correspond to market shocks or fraudulent activity. Tenconi *et al.* (2024) applied a copula-Generalized Linear Model to detect fraudulent bank transfers. By modeling the joint distribution of features (amount, frequency, account histories) via a copula, they accurately estimated the probability of fraud. They found that “*copula-based probability and conditional copula formulas offer promising results, particularly in handling complex dependencies*” in fraud features ¹¹ ¹² . In risk management, copulas have long been used to detect joint tail events (e.g. extreme losses across assets).
- **Climate and Environmental Science:** Climate phenomena are inherently multivariate (temperature, precipitation, pressure, etc.). Copulas are widely used to model joint extremes (e.g. bivariate precipitation–temperature models). For anomaly detection, Bazrafshan *et al.* (2020) used a copula-based method for drought monitoring. They showed that a linear PCA model overestimated anomalies while the copula-based detector was “*not influenced by an anomaly in one or more variables*” ¹⁴ , providing more stable detection of multivariate climate events. Pabaghi *et al.* (2023) used bivariate copulas to model extreme precipitation in arid regions, helping identify unusual rainfall patterns. In hydrology, copulas quantify the joint return period of floods from rainfall and snowmelt ¹⁴ ¹⁶ , effectively flagging compound events (floods) as anomalies in joint climate behavior.
- **Network Traffic and Telecommunications:** Modern networks produce high-dimensional time series of performance counters. Horváth *et al.* (2019) applied a copula-based anomaly scoring method to a large mobile network dataset ⁹ . They modeled the multivariate distribution of hundreds of sensors and detected hidden performance issues that conventional methods missed: “*with the proposed procedure, many hidden issues can be isolated and indicated to the network operator.*” ⁹ . Copula methods also appear in cybersecurity: e.g. copula-based classifiers have been proposed for network intrusion detection, exploiting attribute dependence to spot abnormal traffic patterns.
- **Other Domains:** In manufacturing and IoT sensor networks, copulas detect faults by modeling the normal joint behavior of sensor readings. In geosciences, copulas have been used to detect land cover or ocean anomalies via multivariate environmental indicators. In astronomy and physics, methods like RanBox (2021) scan copula space to discover rare events in high-energy data ¹⁰ .

These case studies illustrate that whenever anomalous behavior arises from unusual **joint patterns** (not just outliers in one variable), copula-based detectors provide a natural solution. By focusing on the dependence structure, they can reveal anomalies that would be hidden to univariate or linear methods.

Conclusion

Copulas offer a principled framework for multivariate anomaly detection in time series. They enable flexible modeling of complex dependencies and tail events, complementing classical and modern techniques. Copula methods provide interpretable anomaly scores (joint p-values) and have been successfully applied in finance, climate science, network monitoring, and beyond. For a researcher, copula-based anomaly detection opens avenues for modeling joint behavior and pooling evidence across series. Future directions include efficient copula estimation for very high dimensions (e.g. regularized vines), dynamic copulas for streaming data, and hybrid methods combining deep learning with copula-based statistical tests.

Sources: Foundational copula theory is given by Sklar’s theorem ³ ; tail-dependence properties by Nagler & Czado ⁵ ; COPOD’s method by Zhao *et al.* ⁷ ; network-anomaly case by Horváth *et al.* ⁹ ; fraud detection example by Tenconi *et al.* ¹¹ ¹² ; multivariate anomaly definitions by Verma *et al.* ² ; comparisons with PCA by Bazrafshan *et al.* ¹⁴ ; and a survey of deep methods ¹⁵ . These and other references are cited above.

1 13 esd.copernicus.org

<https://esd.copernicus.org/preprints/esd-2016-51/esd-2016-51.pdf>

2 Learning Local Patterns of Time Series for Anomaly Detection

<https://www.mdpi.com/2673-4591/39/1/82>

3 columbia.edu

<http://www.columbia.edu/~mh2078/QRM/Copulas.pdf>

4 5 6 tnagler.github.io

<https://tnagler.github.io/vine-arisa.pdf>

7 [2009.09463] COPOD: Copula-Based Outlier Detection

<https://arxiv.org/abs/2009.09463>

8 Using Copula to Model Dependence When Testing Multiple Hypotheses in DNA Microarray Experiments:
A Bayesian Approximation

<https://www.mdpi.com/2227-7390/8/9/1514>

9 [1912.02166] Copula-based anomaly scoring and localization for large-scale, high-dimensional
continuous data

<https://arxiv.org/abs/1912.02166>

10 [2106.05747] RanBox: Anomaly Detection in the Copula Space

<https://arxiv.org/abs/2106.05747>

11 12 Copula-based Approaches for Anomaly Detection: a Case-study in Financial Forensics

<https://re.public.polimi.it/handle/11311/1285650>

14 16 Bivariate Analysis of Extreme Precipitation Using Copula Functions in Arid and Semi-Arid Regions

<https://www.mdpi.com/2073-4433/14/2/275>

15 A Survey of Deep Anomaly Detection in Multivariate Time Series: Taxonomy, Applications, and Directions

<https://www.mdpi.com/1424-8220/25/1/190>